

Notebook 1 Regression

Matthew Naughton

CS 4375.003

Portfolio: Kernel and Ensemble Methods

```
set.seed(1234)
df <- read.csv(file = 'housing.csv')

# install.packages("ggplot2") ##uncomment and run if not installed
# install.packages("e1071")   ##uncomment and run if not installed
# install.packages("MASS")    ##uncomment and run if not installed
library(ggplot2)
library(e1071)
library(MASS)

ocean_factor <- as.factor(df$ocean_proximity)
spec <- c(train=.6, test=.2, validate=.2)
i <- sample(cut(1:nrow(df),
               nrow(df)*cumsum(c(0,spec))), labels=names(spec))
train <- df[i=="train",]
test <- df[i=="test",]
vald <- df[i=="validate",]
#i <- sample(1:nrow(df), nrow(df)*0.75, replace=FALSE)
#train <- df[i,]
#test <- df[-i,]

# Exploring the data:
head(df)

##  longitude latitude housing_median_age total_rooms total_bedrooms
population
## 1  -122.23    37.88           41           880           129
322
## 2  -122.22    37.86           21          7099          1106
2401
## 3  -122.24    37.85           52          1467           190
496
## 4  -122.25    37.85           52          1274           235
558
## 5  -122.25    37.85           52          1627           280
565
## 6  -122.25    37.85           52           919           213
413
##  households median_income median_house_value ocean_proximity
## 1         126         8.3252         452600      NEAR BAY
```

```
## 2      1138      8.3014      358500      NEAR BAY
## 3       177      7.2574      352100      NEAR BAY
## 4       219      5.6431      341300      NEAR BAY
## 5       259      3.8462      342200      NEAR BAY
## 6       193      4.0368      269700      NEAR BAY
```


```
tail(df)
```

```
##      longitude latitude housing_median_age total_rooms total_bedrooms
## 20635   -121.56   39.27           28         2332         395
## 20636   -121.09   39.48           25         1665         374
## 20637   -121.21   39.49           18          697         150
## 20638   -121.22   39.43           17         2254         485
## 20639   -121.32   39.43           18         1860         409
## 20640   -121.24   39.37           16         2785         616
```

```
##      population households median_income median_house_value
ocean_proximity
```

```
## 20635      1041       344       3.7125       116800
INLAND
## 20636       845       330       1.5603       78100
INLAND
## 20637       356       114       2.5568       77100
INLAND
## 20638      1007       433       1.7000       92300
INLAND
## 20639       741       349       1.8672       84700
INLAND
## 20640      1387       530       2.3886       89400
INLAND
```

```
plot(df$median_house_value, df$median_income, xlab="Median House Value",
ylab="Median Income")
```



```
str(df)

## 'data.frame': 20640 obs. of 10 variables:
## $ longitude : num -122 -122 -122 -122 -122 ...
## $ latitude : num 37.9 37.9 37.9 37.9 37.9 ...
## $ housing_median_age: num 41 21 52 52 52 52 52 52 42 52 ...
## $ total_rooms : num 880 7099 1467 1274 1627 ...
## $ total_bedrooms : num 129 1106 190 235 280 ...
## $ population : num 322 2401 496 558 565 ...
## $ households : num 126 1138 177 219 259 ...
## $ median_income : num 8.33 8.3 7.26 5.64 3.85 ...
## $ median_house_value: num 452600 358500 352100 341300 342200 ...
## $ ocean_proximity : chr "NEAR BAY" "NEAR BAY" "NEAR BAY" "NEAR BAY"
...

names(df)

## [1] "longitude" "latitude" "housing_median_age"
## [4] "total_rooms" "total_bedrooms" "population"
## [7] "households" "median_income" "median_house_value"
## [10] "ocean_proximity"

summary(df)

## longitude latitude housing_median_age total_rooms
## Min. :-124.3 Min. :32.54 Min. : 1.00 Min. : 2
## 1st Qu.: -121.8 1st Qu.:33.93 1st Qu.:18.00 1st Qu.: 1448
## Median : -118.5 Median :34.26 Median :29.00 Median : 2127
```

```

## Mean      :-119.6    Mean      :35.63    Mean      :28.64    Mean      : 2636
## 3rd Qu.: -118.0    3rd Qu.:37.71    3rd Qu.:37.00    3rd Qu.: 3148
## Max.      :-114.3    Max.      :41.95    Max.      :52.00    Max.      :39320
##
## total_bedrooms      population      households      median_income
## Min.      : 1.0      Min.      : 3      Min.      : 1.0      Min.      : 0.4999
## 1st Qu.: 296.0      1st Qu.: 787      1st Qu.: 280.0      1st Qu.: 2.5634
## Median : 435.0      Median : 1166      Median : 409.0      Median : 3.5348
## Mean      : 537.9      Mean      : 1425      Mean      : 499.5      Mean      : 3.8707
## 3rd Qu.: 647.0      3rd Qu.: 1725      3rd Qu.: 605.0      3rd Qu.: 4.7432
## Max.      :6445.0      Max.      :35682      Max.      :6082.0      Max.      :15.0001
## NA's      :207
## median_house_value ocean_proximity
## Min.      : 14999      Length:20640
## 1st Qu.:119600      Class :character
## Median :179700      Mode  :character
## Mean      :206856
## 3rd Qu.:264725
## Max.      :500001
##
#range of value and income
range(df$median_house_value)

## [1] 14999 500001


range(df$median_income)

## [1] 0.4999 15.0001


ggplot(data=df)+geom_histogram(mapping = aes(x=median_house_value))

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

```



```
ggplot(df, aes(x=median_house_value, y=median_income))+  
  geom_point()+  
  stat_smooth()  
  
## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```



```
#Correlation:
cor(df$median_house_value, df$median_income)

## [1] 0.6880752

# Linear regression
lm1 <- lm(median_house_value~median_income, data=train)
pred <- predict(lm1, newdata=test)
cor_lm1 <- cor(pred, test$median_house_value)
mse_lm1 <- mean((pred - test$median_house_value)^2)

# SVM linear kernel
svm1 <- svm(median_house_value~., data=train, kernel="linear", cost=10,
scale=TRUE)
summary(svm1)

##
## Call:
## svm(formula = median_house_value ~ ., data = train, kernel = "linear",
##      cost = 10, scale = TRUE)
##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: linear
##      cost:   10
##   gamma:    0.07692308
```

[illegible]

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```



```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of  
shorter  
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
```

[illegible]

```
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
```

```
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
```

```
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
```

```
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
```

```
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
```

```
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
```

```
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
```

```
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
```

```
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
```

```
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
```

```
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
```

```
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
```

```
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
## object length
```

```
## Warning in pred - true.y: longer object length is not a multiple of
shorter
## object length
```

```
summary(tune_svm1)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.001
##
## - best performance: 13516180713
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-03 13516180713 3167630515
## 2 1e-02 14618991604 3668035847
## 3 1e-01 14755060141 3721498741
## 4 1e+00 14768185264 3722444558
## 5 5e+00 14772729803 3725927367
## 6 1e+01 14773799700 3726770641
```

```
pred <- predict(tune_svm1$best.model, newdata=test)
cor_svm1_tune <- cor(pred, test$median_house_value[(1:length(pred))])
mse_svm1_tune <- mean((pred - test$median_house_value[(1:length(pred))])^2)
```

```
# SVM polynomial kernel
```

```
svm2 <- svm(median_house_value~., data=train, kernel="polynomial", cost=10,
scale=TRUE)
summary(svm2)
```

```
##
## Call:
## svm(formula = median_house_value ~ ., data = train, kernel = "polynomial",
##   cost = 10, scale = TRUE)
##
##
```

```

## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: polynomial
##     cost: 10
##     degree: 3
##     gamma: 0.07692308
##     coef.0: 0
##     epsilon: 0.1
##
##
## Number of Support Vectors: 9260

# Evaluate the polynomial kernel

pred <- predict(svm2, newdata=test)
cor_svm2 <- cor(pred, test$median_house_value[(1:length(pred))])
mse_svm2 <- mean((pred - test$median_house_value[(1:length(pred))])^2)

#table(pred2, test$ocean_proximity[(1:length(pred2))])
#mean(pred2==test$ocean_proximity[(1:length(pred2))])
#plot(svm2, test, median_income ~ median_house_value)

# SVM radial kernel

svm3 <- svm(median_house_value~., data=train, kernel="radial", cost=10,
gamma=1, scale=TRUE)
summary(svm3)

##
## Call:
## svm(formula = median_house_value ~ ., data = train, kernel = "radial",
##     cost = 10, gamma = 1, scale = TRUE)
##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: radial
##     cost: 10
##     gamma: 1
##     epsilon: 0.1
##
##
## Number of Support Vectors: 9073

# Evaluate radial kernel

pred <- predict(svm3, newdata=test)
cor_svm3 <- cor(pred, test$median_house_value[(1:length(pred))])
mse_svm3 <- mean((pred - test$median_house_value[(1:length(pred))])^2)

#table(pred, test$median_house_value[(1:length(pred))])

```

```

#mean(pred==test$median_house_value[(1:length(pred))])
#plot(svm3, test, median_house_value ~ median_income)

# Tuning the hyperparameters
tune.out <- tune(svm, median_house_value~., data=valid, kernel="radial",
                ranges=list(cost=c(0.1,1,10),
                             gamma=c(0.5,1,2,3,4)))
summary(tune.out)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##     1    0.5
##
## - best performance: 3643851125
##
## - Detailed performance results:
##   cost gamma      error dispersion
## 1  0.1   0.5  5556095639  702450336
## 2  1.0   0.5  3643851125  401025212
## 3 10.0   0.5  3710001847  439752018
## 4  0.1   1.0  7288940427  921340109
## 5  1.0   1.0  4338342266  494865720
## 6 10.0   1.0  4515736157  390752698
## 7  0.1   2.0  9371781168 1053655868
## 8  1.0   2.0  5745424123  711741863
## 9 10.0   2.0  5692900032  474132066
##10  0.1   3.0 10615120576 1088078083
##11  1.0   3.0  6778754620  845277922
##12 10.0   3.0  6555410666  618531944
##13  0.1   4.0 11440269200 1098435926
##14  1.0   4.0  7614682842  938718401
##15 10.0   4.0  7295446133  734854825

# Radial kernel with various cost and gamma values
svm4 <- svm(median_house_value~., data=train, kernel = "radial", cost=100,
            gamma=0.5, scale=TRUE)
summary(svm4)

##
## Call:
## svm(formula = median_house_value ~ ., data = train, kernel = "radial",
##     cost = 100, gamma = 0.5, scale = TRUE)
##
##
## Parameters:
##   SVM-Type:  eps-regression

```

```

## SVM-Kernel: radial
##      cost: 100
##      gamma: 0.5
##      epsilon: 0.1
##
## Number of Support Vectors: 9020

# Evaluate Radial kernel with various cost/gamma values
pred <- predict(svm4, newdata=test)
cor_svm4 <- cor(pred, test$median_house_value[(1:length(pred))])
mse_svm4 <- mean((pred - test$median_house_value[(1:length(pred))])^2)

#table(pred, test$ocean_proximity[(1:length(pred))])
#mean(pred==test$ocean_proximity[(1:length(pred))])
#plot(svm4, test, median_income ~ median_house_value)

# correlation and mse for linear model
cor_lm1

## [1] 0.6980403

mse_lm1

## [1] 7052106439

# svm1 through svm4 correlation and mse
cor_svm1

## [1] 0.3748647

mse_svm1

## [1] 14472254737

cor_svm2

## [1] 0.114215

mse_svm2

## [1] 98206193678

cor_svm3

## [1] 0.4225706

mse_svm3

## [1] 14317898140

cor_svm4

## [1] 0.4202958

```

```
mse_svm4
```

```
## [1] 15123977772
```

#Overall, the linear model had the best correlation value, I believe this is because the data is more linear and therefore attempts to polynomialize the data would make it worse.